

R: Spatial regression

Roger Bivand

Thursday, 5 September 2019, 15:12-15:36

Required current contributed CRAN packages:

I am running R 3.6.1, with recent `update.packages()`.

```
needed <- c("MatrixModels", "lme4", "spatialreg", "spdep", "sf", "sp", "HSAR")
```

Beijing data set

Lower level point support data and upper level district boundaries (polygon support)

```
library(HSAR)
library(sp)
data(landSPDF)
data(landprice)
data(Beijingdistricts)
```

Convert to `sf` class and merge data with point geometries

```
library(sf)
```

```
## Linking to GEOS 3.7.2, GDAL 3.0.1, PROJ 6.2.0
```

```
land_sf <- st_as_sf(landSPDF)
landprice_sf <- merge(land_sf, landprice, by="obs")
(landprice_sf <- landprice_sf[order(landprice_sf$district.id.x),])
```

```
## Simple feature collection with 1117 features and 12 fields
```

```
## geometry type: POINT
```

```
## dimension: XY
```

```
## bbox: xmin: 428553.1 ymin: 4406815 xmax: 463693.2 ymax: 4440423
```

```
## epsg (SRID): NA
```

```
## proj4string: +proj=tmerc +lat_0=0 +lon_0=117 +k=1 +x_0=500000 +y_0=0 +ellps=krass +units=m +no_defs
```

```
## First 10 features:
```

```
##      obs district.id.x lnprice  lnarea  lndcbd dsubway  dpark  dele
## 187 189              3 5.57430 10.27820 9.94866 6.83023 7.06579 6.81916
## 188 190              3 7.16382 11.58780 9.93534 7.14334 6.78243 6.67827
## 700 968              3 7.61282  8.94551 9.91779 7.64360 6.84364 4.60356
## 701 969              3 6.81564  5.81928 9.91940 7.64640 6.88254 4.10025
## 702 970              3 6.93528  7.71869 9.91752 7.65810 6.86760 4.53460
## 709 992              3 7.45757  9.20029 9.84785 7.78904 6.95662 7.05138
## 710 993              3 7.12569  7.97788 9.84388 7.81991 7.00792 7.11267
## 711 994              3 7.48522  7.78634 9.84203 7.83398 7.03089 7.13981
## 717 1001             3 5.87349 10.70910 9.95534 7.89121 7.11019 5.67984
## 181 183              5 6.79302  6.39403 9.92025 6.76006 6.23524 6.10494
```

```
##      popden crimerate district.id.y year      geometry
## 187 0.548966 10.75110              3   1 POINT (430237 4422804)
## 188 0.548966 10.75110              3   1 POINT (430547.1 4423001)
## 700 0.548966 10.75110              3   0 POINT (431029.4 4423667)
## 701 0.548966 10.75110              3   0 POINT (431001.8 4423695)
```

```
## 702 0.548966 10.75110      3  0 POINT (431040.9 4423697)
## 709 0.548966 10.75110      3  0 POINT (432164 4422080)
## 710 0.548966 10.75110      3  0 POINT (432239.7 4422081)
## 711 0.548966 10.75110      3  0 POINT (432275 4422082)
## 717 0.548966 10.75110      3  0 POINT (430436.8 4424601)
## 181 1.407250  2.25832      5  1 POINT (430606.5 4420186)
```

Check that the input IDs match and that the data are correctly ordered

```
all.equal(landprice_sf$district.id.x, landprice_sf$district.id.y)
```

```
## [1] TRUE
```

Create the original 1.5 km distance threshold spatial weights, with a few no-neighbour observations (so set zero policy option)

```
library(spatialreg)
```

```
## Loading required package: spData
```

```
## Loading required package: Matrix
```

```
## Registered S3 methods overwritten by 'spatialreg':
```

```
## method          from
## residuals.stsls  spdep
## deviance.stsls   spdep
## coef.stsls       spdep
## print.stsls      spdep
## summary.stsls    spdep
## print.summary.stsls spdep
## residuals.gmsar  spdep
## deviance.gmsar   spdep
## coef.gmsar       spdep
## fitted.gmsar     spdep
## print.gmsar      spdep
## summary.gmsar    spdep
## print.summary.gmsar spdep
## print.lagmess    spdep
## summary.lagmess   spdep
## print.summary.lagmess spdep
## residuals.lagmess spdep
## deviance.lagmess  spdep
## coef.lagmess      spdep
## fitted.lagmess    spdep
## logLik.lagmess    spdep
## fitted.SFResult   spdep
## print.SFResult    spdep
## fitted.ME_res     spdep
## print.ME_res      spdep
## print.lagImpact   spdep
## plot.lagImpact    spdep
## summary.lagImpact  spdep
## HPDinterval.lagImpact spdep
## print.summary.lagImpact spdep
## print.sarlm       spdep
## summary.sarlm     spdep
## residuals.sarlm   spdep
```

```

## deviance.sarlm          spdep
## coef.sarlm              spdep
## vcov.sarlm              spdep
## fitted.sarlm            spdep
## logLik.sarlm            spdep
## anova.sarlm             spdep
## predict.sarlm           spdep
## print.summary.sarlm     spdep
## print.sarlm.pred        spdep
## as.data.frame.sarlm.pred spdep
## residuals.spautolm      spdep
## deviance.spautolm       spdep
## coef.spautolm           spdep
## fitted.spautolm         spdep
## print.spautolm          spdep
## summary.spautolm        spdep
## logLik.spautolm         spdep
## print.summary.spautolm  spdep
## print.WXImpact          spdep
## summary.WXImpact        spdep
## print.summary.WXImpact  spdep
## predict.SLX             spdep


dnb1.5 <- spdep::dnearneigh(landprice_sf, 0, 1500, row.names=as.character(landprice_sf$obs))
dnb1.5

## Neighbour list object:
## Number of regions: 1117
## Number of nonzero links: 18798
## Percentage nonzero weights: 1.506625
## Average number of links: 16.82901
## 7 regions with no links:
## 517 53 292 1764 33 1785 1126

dists <- spdep::nbdists(dnb1.5, st_geometry(landprice_sf))
edists <- lapply(dists, function(x) exp(-(x/1000)^2)/(1.5^2))
ozpo <- spdep::set.ZeroPolicyOption(TRUE)
oo <- set.ZeroPolicyOption(TRUE)
lw <- spdep::nb2listw(dnb1.5, glist=edists, style="W")
hist(spdep::card(dnb1.5))

```

Histogram of `spdep::card(dnb1.5)`



Reconstruct the input data for R formula use (do not log in advance, do use factors for categorical variables to permit automatic generation of dummies)

```
landprice_sf$fyyear <- factor(landprice_sf$year + 2003)
landprice_sf$price <- exp(landprice_sf$lnprice)
landprice_sf$area <- exp(landprice_sf$lnarea)
landprice_sf$Dcbd <- exp(landprice_sf$lncbd)
landprice_sf$Dsubway <- exp(landprice_sf$dsubway)
landprice_sf$Dpark <- exp(landprice_sf$dpark)
landprice_sf$Dele <- exp(landprice_sf$ddele)
landprice_sf$f_district.id <- factor(landprice_sf$district.id.x)
(t1 <- table(table(landprice_sf$f_district.id)))
```

```
##
##  1  2  3  4  5  6  7  8  9 10 11 12 14 15 16 17 18 19 20 21 23 25 26 27 28
##  7  8 10  4  5 10  8 10  8  4  6  2  4  3  1  2  5  2  1  1  1  1  1  1  2
## 31 32 33 52
##  1  1  1  1
```

Some covariates are observed at the district level rather than the land parcel level

```
supply(as.data.frame(landprice_sf[, c("price", "area", "Dcbd", "Dele", "Dpark", "Dsubway", "crimerate",
## price area Dcbd Dele Dpark Dsubway crimerate
## 1062 1098 1114 1116 1114 1117 105
## popden geometry
## 111 1117
```

Check the matching of district IDs and counts of land parcels in districts

```
Beijingdistricts$id1 <- Beijingdistricts$id+1
all.equal(unique(landprice_sf$district.id.x), Beijingdistricts$id1)
```

```
## [1] TRUE

(Beijingdistricts_sf <- st_as_sf(Beijingdistricts))

## Simple feature collection with 111 features and 2 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:           xmin: 426987.3 ymin: 4403559 xmax: 467920.9 ymax: 4443287
## epsg (SRID):    NA
## proj4string:     +proj=tmerc +lat_0=0 +lon_0=117 +k=1 +x_0=500000 +y_0=0 +ellps=krass +units=m +no_de
## First 10 features:
##   id id1 geometry
## 0  2   3 MULTIPOLYGON (((428183.1 44...
## 1  4   5 MULTIPOLYGON (((432472.2 44...
## 2  6   7 MULTIPOLYGON (((432446.1 44...
## 3  7   8 MULTIPOLYGON (((433534.5 44...
## 4  8   9 MULTIPOLYGON (((443807.9 44...
## 5  9  10 MULTIPOLYGON (((444461 4420...
## 6 10  11 MULTIPOLYGON (((447530.6 44...
## 7 11  12 MULTIPOLYGON (((443849.6 44...
## 8 12  13 MULTIPOLYGON (((446810 4417...
## 9 13  14 MULTIPOLYGON (((445954.8 44...

Beijingdistricts_sf$counts <- sapply(st_contains(Beijingdistricts_sf, landprice_sf), length)
```

Check point counts by district from input data and topological points in polygon counts

```
t2 <- table(Beijingdistricts_sf$counts)
all.equal(t1, t2)
```

```
## [1] TRUE
```

Basic formula object from the original paper and examples in package, fit and display OLS model (note that fyear is split into dummies with 2003 in the intercept)

```
form <- log(price) ~ log(area) + log(Dcbd) + log(Dele) + log(Dpark) + log(Dsubway) +
  crimerate + popden + fyear
OLS <- lm(form, data=landprice_sf)
summary(OLS)
```

```
##
## Call:
## lm(formula = form, data = landprice_sf)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5915 -0.5752 -0.0496  0.5206  3.4042
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  14.153917   0.370103  38.243  < 2e-16 ***
## log(area)    -0.008253   0.018675  -0.442  0.65863
## log(Dcbd)    -0.250601   0.047752  -5.248  1.84e-07 ***
## log(Dele)    -0.085528   0.032308  -2.647  0.00823 **
## log(Dpark)   -0.284372   0.046046  -6.176  9.24e-10 ***
## log(Dsubway) -0.245748   0.034755  -7.071  2.73e-12 ***
## crimerate     0.007668   0.004458   1.720  0.08575 .
```

```
## popden      0.032827  0.010304  3.186  0.00148 **
## fyear2004   -0.164503  0.058380 -2.818  0.00492 **
## fyear2005    0.017635  0.124986  0.141  0.88782
## fyear2006   -0.120314  0.107209 -1.122  0.26201
## fyear2007    0.551384  0.117431  4.695 3.00e-06 ***
## fyear2008    0.396172  0.129571  3.058  0.00229 **
## fyear2009    2.113691  0.228688  9.243 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8335 on 1103 degrees of freedom
## Multiple R-squared:  0.3524, Adjusted R-squared:  0.3448
## F-statistic: 46.17 on 13 and 1103 DF, p-value: < 2.2e-16
```

Are the residuals spatially autocorrelated?

```
spdep::lm.morantest(OLS, listw=lw)
```

```
##
## Global Moran I for regression residuals
##
## data:
## model: lm(formula = form, data = landprice_sf)
## weights: lw
##
## Moran I statistic standard deviate = 15.1, p-value < 2.2e-16
## alternative hypothesis: greater
## sample estimates:
## Observed Moran I      Expectation      Variance
## 0.1944768641      -0.0054494313      0.0001753028
```

What do the robust LM tests say?

```
spdep::lm.LMtests(OLS, listw=lw, test=c("RLMerr", "RLMlag"))
```

```
##
## Lagrange multiplier diagnostics for spatial dependence
##
## data:
## model: lm(formula = form, data = landprice_sf)
## weights: lw
##
## RLMerr = 118.22, df = 1, p-value < 2.2e-16
##
## Lagrange multiplier diagnostics for spatial dependence
##
## data:
## model: lm(formula = form, data = landprice_sf)
## weights: lw
##
## RLMlag = 3.7725, df = 1, p-value = 0.0521
```

How do we do with a linear model including selected spatially lagged covariates?

```
SLX <- lmSLX(form, data=landprice_sf, listw=lw, Durbin=~ log(area) + log(Dcbd) + log(Dele) + log(Dpark),
summary(impacts(SLX))
```

```

## Impact measures (SLX, estimable, n-k):
##           Direct      Indirect      Total
## log(area)   -0.028250900  0.15061221  0.122361308
## log(Dcbd)   -0.575325229  0.34199016 -0.233335072
## log(Dele)    0.001927612 -0.15153784 -0.149610227
## log(Dpark)  -0.066376027 -0.29670998 -0.363086011
## log(Dsubway) -0.180677830 -0.09171154 -0.272389368
## crimerate    0.002228943  0.00697053  0.009199473
## popden       0.003782010  0.04501340  0.048795406
## fyear2004   -0.163634606      NA -0.163634606
## fyear2005    0.006112745      NA  0.006112745
## fyear2006   -0.120821815      NA -0.120821815
## fyear2007    0.619330756      NA  0.619330756
## fyear2008    0.444819013      NA  0.444819013
## fyear2009    2.186988035      NA  2.186988035
## =====
## Standard errors:
##           Direct      Indirect      Total
## log(area)   0.019588590  0.03882315  0.037437930
## log(Dcbd)   0.121875390  0.13150000  0.052242219
## log(Dele)   0.051081439  0.06620277  0.043070631
## log(Dpark)  0.107848818  0.12908527  0.057968077
## log(Dsubway) 0.059325583  0.07722793  0.046759734
## crimerate   0.009157083  0.01134039  0.005658334
## popden      0.017573848  0.02455712  0.014701241
## fyear2004   0.058433277      NA  0.058433277
## fyear2005   0.124199473      NA  0.124199473
## fyear2006   0.106956715      NA  0.106956715
## fyear2007   0.117305346      NA  0.117305346
## fyear2008   0.128990353      NA  0.128990353
## fyear2009   0.227449165      NA  0.227449165
## =====
## Z-values:
##           Direct      Indirect      Total
## log(area)   -1.44221201  3.8794431  3.26837801
## log(Dcbd)   -4.72060215  2.6006857 -4.46640817
## log(Dele)    0.03773605 -2.2889955 -3.47360195
## log(Dpark)  -0.61545437 -2.2985581 -6.26355107
## log(Dsubway) -3.04552976 -1.1875436 -5.82529759
## crimerate    0.24341193  0.6146642  1.62582705
## popden       0.21520671  1.8330078  3.31913527
## fyear2004   -2.80036674      NA -2.80036674
## fyear2005    0.04921716      NA  0.04921716
## fyear2006   -1.12963282      NA -1.12963282
## fyear2007    5.27964648      NA  5.27964648
## fyear2008    3.44846729      NA  3.44846729
## fyear2009    9.61528276      NA  9.61528276
##
## p-values:
##           Direct      Indirect      Total
## log(area)   0.14924257  0.0001047  0.00108166
## log(Dcbd)   2.3515e-06  0.0093038  7.9544e-06
## log(Dele)   0.96989813  0.0220796  0.00051352
## log(Dpark)  0.53825469  0.0215300  3.7631e-10

```

```
## log(Dsubway) 0.00232271 0.2350133 5.7011e-09
## crimerate   0.80768630 0.5387765 0.10398645
## popden      0.82960616 0.0668014 0.00090297
## fyear2004   0.00510446 NA      0.00510446
## fyear2005   0.96074624 NA      0.96074624
## fyear2006   0.25863098 NA      0.25863098
## fyear2007   1.2943e-07 NA      1.2943e-07
## fyear2008   0.00056378 NA      0.00056378
## fyear2009   < 2.22e-16 NA      < 2.22e-16
```

And the spatial residual autocorrelation?

```
spdep::lm.morantest(SLX, listw=lw)
```

```
##
## Global Moran I for regression residuals
##
## data:
## model: lm(formula = formula(paste("y ~ ", paste(colnames(x)[-1],
## collapse = "+"))), data = as.data.frame(x), weights = weights)
## weights: lw
##
## Moran I statistic standard deviate = 14.802, p-value < 2.2e-16
## alternative hypothesis: greater
## sample estimates:
## Observed Moran I      Expectation      Variance
##      0.1885391149    -0.0069886190      0.0001745023
```

And robust LM tests?

```
spdep::lm.LMtests(SLX, listw=lw, test=c("RLMerr", "RLMlag"))
```

```
##
## Lagrange multiplier diagnostics for spatial dependence
##
## data:
## model: lm(formula = formula(paste("y ~ ", paste(colnames(x)[-1],
## collapse = "+"))), data = as.data.frame(x), weights = weights)
## weights: lw
##
## RLMerr = 39.985, df = 1, p-value = 2.56e-10
##
##
## Lagrange multiplier diagnostics for spatial dependence
##
## data:
## model: lm(formula = formula(paste("y ~ ", paste(colnames(x)[-1],
## collapse = "+"))), data = as.data.frame(x), weights = weights)
## weights: lw
##
## RLMlag = 0.17014, df = 1, p-value = 0.68
```

So let's fit a spatial Durbin error model, with the same selection of spatially lagged covariates

```
e <- eigenw(lw)
SDEM <- errorsarlm(form, data=landprice_sf, listw=lw, Durbin= ~ log(area) + log(Dcbd) + log(Dele) + log
summary(impacts(SDEM))
```



```

## Impact measures (SDEM, estimable, n):
##           Direct      Indirect      Total
## log(area)   -0.021085385  0.133094132  0.112008747
## log(Dcbd)   -0.572352697  0.317284147 -0.255068551
## log(Dele)   -0.005661913 -0.114782232 -0.120444145
## log(Dpark)  -0.074378191 -0.273739346 -0.348117537
## log(Dsubway) -0.164798714 -0.090024131 -0.254822845
## crimerate    0.006411748 -0.001121658  0.005290091
## popden       0.004983237  0.045377975  0.050361213
## fyear2004   -0.213875338          NA -0.213875338
## fyear2005   -0.074548268          NA -0.074548268
## fyear2006   -0.136022976          NA -0.136022976
## fyear2007    0.713777760          NA  0.713777760
## fyear2008    0.488689750          NA  0.488689750
## fyear2009    2.167398316          NA  2.167398316
## =====
## Standard errors:
##           Direct      Indirect      Total
## log(area)   0.018070856  0.05086004  0.055146985
## log(Dcbd)   0.124751995  0.14355527  0.091948225
## log(Dele)   0.045634066  0.08014074  0.072364276
## log(Dpark)  0.101149512  0.14212885  0.094724658
## log(Dsubway) 0.054362580  0.09395633  0.078343698
## crimerate   0.008305157  0.01274632  0.009474872
## popden      0.015899732  0.02876860  0.023482266
## fyear2004   0.055676729          NA  0.055676729
## fyear2005   0.116369514          NA  0.116369514
## fyear2006   0.100092117          NA  0.100092117
## fyear2007   0.115930158          NA  0.115930158
## fyear2008   0.122087311          NA  0.122087311
## fyear2009   0.205699716          NA  0.205699716
## =====
## Z-values:
##           Direct      Indirect      Total
## log(area)   -1.1668171  2.61687048  2.0310947
## log(Dcbd)   -4.5879242  2.21018802 -2.7740454
## log(Dele)   -0.1240721 -1.43225824 -1.6644144
## log(Dpark)  -0.7353292 -1.92599427 -3.6750467
## log(Dsubway) -3.0314734 -0.95814867 -3.2526272
## crimerate    0.7720201 -0.08799852  0.5583285
## popden       0.3134164  1.57734403  2.1446488
## fyear2004   -3.8413776          NA -3.8413776
## fyear2005   -0.6406168          NA -0.6406168
## fyear2006   -1.3589779          NA -1.3589779
## fyear2007    6.1569636          NA  6.1569636
## fyear2008    4.0027890          NA  4.0027890
## fyear2009   10.5367103          NA 10.5367103
##
## p-values:
##           Direct      Indirect Total
## log(area)   0.24328423  0.008874  0.04224539
## log(Dcbd)   4.4768e-06  0.027092  0.00553640
## log(Dele)   0.90125821  0.152070  0.09602964
## log(Dpark)  0.46213902  0.054105  0.00023781

```

```
## log(Dsubway) 0.00243363 0.337988 0.00114343
## crimerate   0.44010252 0.929878 0.57662011
## popden      0.75396431 0.114716 0.03198094
## fyear2004   0.00012235 NA         0.00012235
## fyear2005   0.52177167 NA         0.52177167
## fyear2006   0.17415359 NA         0.17415359
## fyear2007   7.4153e-10 NA         7.4153e-10
## fyear2008   6.2600e-05 NA         6.2600e-05
## fyear2009   < 2.22e-16 NA         < 2.22e-16
```

The likelihood ratio test shows that the SDEM model fits much better than the SLX model

```
LR1.sarlm(SDEM)
```

```
##
## Likelihood Ratio diagnostics for spatial dependence
##
## data:
## Likelihood ratio = 117.78, df = 1, p-value < 2.2e-16
## sample estimates:
## Log likelihood of spatial error model
##                               -1299.824
##      Log likelihood of OLS fit y
##                               -1358.713
```

The Hausman test is perhaps significant, suggestion that the non-spatial coefficients shift somewhat between the SLX model and the SDEM model

```
Hausman.test(SDEM)
```

```
##
## Spatial Hausman test (asymptotic)
##
## data: NULL
## Hausman test = 41.781, df = 21, p-value = 0.004482
```

But ...

is this the end of the story? Reach out to very general mixed model IID random effects at the district level (fixed effects would give 111 dummies); here without spatially lagged covariates

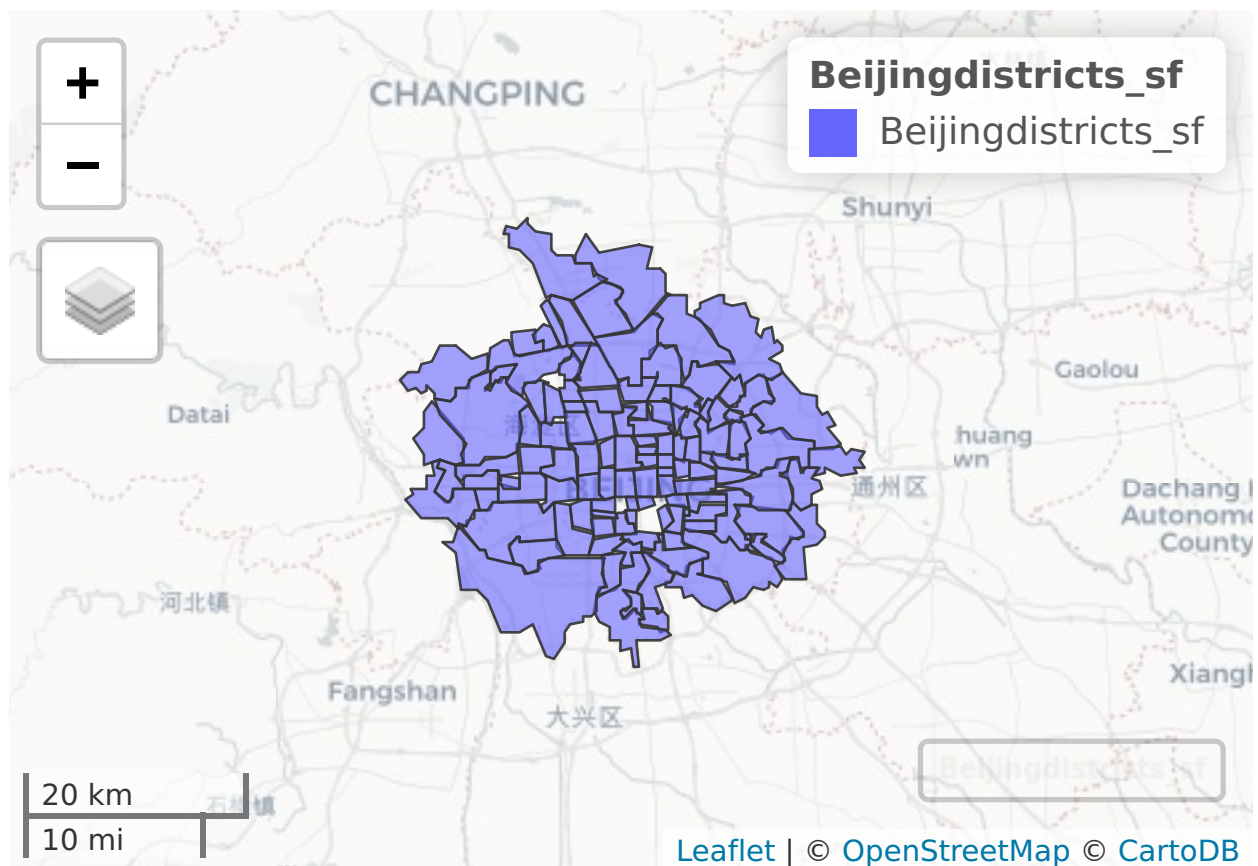
```
library(lme4)
mlm_1 <- lmer(update(form, . ~ . + (1 | f_district.id)), data=landprice_sf, REML=FALSE)
Beijingdistricts_sf$mlm_re <- ranef(mlm_1)[[1]][,1]
```

The **HSAR** model gives a spatial error model at the district level, defining a sparse matrix **Delta** assigning parcels to districts


```
library(Matrix)
suppressMessages(library(MatrixModels))
Delta <- as(model.Matrix(~ -1 + f_district.id, data=landprice_sf, sparse=TRUE), "dgCMatrix")
```

There are gaps in the land parcel and district coverage

```
library(mapview)
mapview(Beijingdistricts_sf)
```



```
opar <- par(bg="lightgreen")
plot(Beijingdistricts_sf[, "counts"])
```



```
par(opar)
```


Construct the spatial weights for the districts

```
nb_M <- spdep::poly2nb(Beijingdistricts, queen=FALSE, row.names=as.character(Beijingdistricts$id1))
M <- as(spdep::nb2listw(nb_M, style="B"), "CsparseMatrix")
dim(M)
```

```
## [1] 111 111
```

```
hist(spdep::card(nb_M))
```

Histogram of `spdep::card(nb_M)`




Using the `M` sparse spatial weights matrix, fit a model with district level simultaneous error autoregression; without spatially lagged covariates

```
m_hsar <- hsar(form, data=landprice_sf, W=NULL, M=M, Delta=Delta, burnin=500, Nsim=5000, thinning=1)
Beijingdistricts_sf$hsar_re <- m_hsar$Mus[1,]
```

The IID and SAR random effects are rather similar


```
plot(Beijingdistricts_sf[, "mlm_re"])
```

mlm_re



```
plot(Beijingdistricts_sf[, "hsar_re"])
```

hsar_re



We do not have tests for residual autocorrelation for these fitted multilevel models, so (speculatively) let's

copy out the district level random effects to the parcels, checking first for matching

```
o <- match(landprice_sf$district.id.x, Beijingdistricts_sf$id1)
landprice_sf$id1 <- Beijingdistricts_sf$id1[o]
all.equal(landprice_sf$district.id.x, landprice_sf$id1)
```

```
## [1] TRUE
```

```
landprice_sf$mlm_re <- Beijingdistricts_sf$mlm_re[o]
landprice_sf$hsar_re <- Beijingdistricts_sf$hsar_re[o]
```

Now we can refit the SLX model but including the district level IID random effect, and test the residual autocorrelation

```
spdep::lm.morantest(lmSLX(update(form, . ~ . + mlm_re), data=landprice_sf, listw=lw, Durbin= ~ log(area
```

```
##
## Global Moran I for regression residuals
##
## data:
## model: lm(formula = formula(paste("y ~ ", paste(colnames(x)[-1],
## collapse = "+"))), data = as.data.frame(x), weights = weights)
## weights: lw
##
## Moran I statistic standard deviate = 2.7727, p-value = 0.002779
## alternative hypothesis: greater
## sample estimates:
## Observed Moran I      Expectation      Variance
##      0.028860257      -0.007654903      0.000173432
```

and for the spatially structured random effect

```
spdep::lm.morantest(lmSLX(update(form, . ~ . + hsar_re), data=landprice_sf, listw=lw, Durbin= ~ log(area
```

```
##
## Global Moran I for regression residuals
##
## data:
## model: lm(formula = formula(paste("y ~ ", paste(colnames(x)[-1],
## collapse = "+"))), data = as.data.frame(x), weights = weights)
## weights: lw
##
## Moran I statistic standard deviate = 3.3167, p-value = 0.0004554
## alternative hypothesis: greater
## sample estimates:
## Observed Moran I      Expectation      Variance
##      0.0359635684      -0.0077079227      0.0001733743
```

Next fit SDEM models with the IID random effect, and it turns out that the SLX model does almost as well, so maybe most of the residual autocorrelation was at the district level rather than the parcel level?

```
SDEM1 <- errorsarlm(update(form, . ~ . + mlm_re), data=landprice_sf, listw=lw, Durbin= ~ log(area) + log
```

```
LR1.sarlm(SDEM1)
```

```
##
## Likelihood Ratio diagnostics for spatial dependence
##
## data:
```

```
## Likelihood ratio = 4.3635, df = 1, p-value = 0.03672
## sample estimates:
## Log likelihood of spatial error model
##               -1224.315
##      Log likelihood of OLS fit y
##               -1226.496
```

The SSRE doesn't do as well (perhaps because it oversmooths the districts)

```
SDEM2 <- errorsarlm(update(form, . ~ . + hsar_re), data=landprice_sf, listw=lw, Durbin= ~ log(area) + 1
LR1.sarlm(SDEM2)
```

```
##
## Likelihood Ratio diagnostics for spatial dependence
##
## data:
## Likelihood ratio = 6.6411, df = 1, p-value = 0.009965
## sample estimates:
## Log likelihood of spatial error model
##               -1229.587
##      Log likelihood of OLS fit y
##               -1232.908
```

Isn't spatial regression fun!!